PolyBench C: speedup over backend-native baseline

GPU: native CUDA baseline • CPU: native C baseline • retained minimum; protocol differences footnoted • higher is better

Speedup (%) = (backend-native minimum / candidate minimum − 1) × 100. Bars are capped at +500%; labels report the actual value.

× = no comparable measured result. † = PolyBench-ACC exists, but uses STANDARD rather than LARGE and is not assigned a speedup.

B300 panels use precision-matched native PolyBenchGPU CUDA as the 0% baseline. The x86 panel uses native FP32/LARGE C on the CPU as its baseline.

KernelFaRer 2mm and 3mm use corrected loop placement; complete outputs pass and the executed paths contain two and three non-nested CBLAS calls, respectively.

† Native PolyBenchGPU FP32 and Polygeist FP32 ran on different physical B300 workers. Their ratios are secondary cross-worker evidence, not strict same-device speedups.

The x86 KernelFaRer, Pluto, and Polly series use one warmup plus ten measurements; x86 Polygeist uses five warmups plus five measurements. They are secondary comparisons.